

Experiment 8: Well-Documented Poets Only

114 Surahs as Separate Authors — Results

Restricted the poet pool to only the 35 poets (of 260) whose total surviving output exceeds 286 verses — the length of the largest surah — then reclustered all 114 surahs (each as its own separate author) against just this well-documented group.

Headline Result

Zero surahs shared a cluster with a poet. Every single one of the 114 surahs grouped only with other surahs. This is the clearest result of any experiment run so far.

Before / After Comparison

Metric	All 260 poets (Exp 2)	35 well-documented poets (Exp 8)
Surahs sharing a cluster with a poet	30 / 114 (26%)	0 / 114 (0%)
Surahs as outliers	1 / 114	3 / 114
Silhouette score	0.469	0.679 (much cleaner)
Total clusters	12	9




Figure 1. Mixing percentage and clustering quality, before vs. after restricting to well-documented poets.

Cluster Composition

With well-documented poets isolated, the two groups separate completely: one cluster is 100% poets (35 of them, all in a single cluster together), and every other cluster is 100% surahs. No cluster contains both types.

Cluster	Poets	Surahs
Cluster 0	35	0
Cluster 1	0	22
Cluster 2	0	10
Cluster 3	0	26
Cluster 4	0	8
Cluster 5	0	22
Cluster 6	0	7
Cluster 7	0	6
Cluster 8	0	10
Outliers	0	3




Figure 2. UMAP projection. Blue circles = the 35 well-documented poets (one tight group, top-left). Triangles = all 114 surahs, colored by their own sub-cluster, forming a completely separate arc.

What This Means for the Sample-Size Question

This is now the second independent test pointing the same direction. Earlier, restricting to MINOR poets (output similar in size to one surah) *increased* the overlap with the Quran, from 26% to 39%. Here, restricting to the opposite group — well-documented poets with substantial surviving output — *eliminated* the overlap entirely, down to 0%, while also producing a much cleaner overall clustering (silhouette 0.679 vs. 0.469).

Both results move in the direction the sample-size explanation predicts: poets with less surviving text produce less distinctive embeddings that sit closer to everything, including the Quran; poets with more text produce clearly distinctive embeddings that separate cleanly. That's a consistent, two-sided result, not a coincidence in one direction only.

Note: One Comparison Still Open

This completes the "114 surahs as separate authors" side of the comparison grid (Exp 2 vs. Exp 8). The "Quran as one author" side has Experiment 1 confirmed (261, Quran shared a cluster with 19 poets), but Experiment 7's result (well-documented poets + Quran as one) hasn't been reported back yet — worth including for the complete picture.

What This Does and Does Not Prove

This still measures writing-style similarity, not authorship. A cleaner separation between well-documented poets and the Quran describes a pattern in the embeddings — consistent with a sample-size effect as explained above — not evidence about who wrote what. The result is more methodologically clean than earlier runs, but the interpretive limits are the same as every report before it.