

Final Summary: Poets + Quran Clustering

Complete 2x2 Comparison — Experiments 1, 2, 7, 8

Four experiments, testing two design choices against each other: how the Quran is represented (as one combined entity, or as 114 separate surahs), crossed with which poets are included (all 260, or only the 35 with substantial surviving output).

The Complete Grid

	All 260 poets	35 well-documented poets only
Quran as ONE entity	Exp 1: 19 poets share its cluster (top similarity 0.976)	Exp 7: 9 poets share its cluster (top similarity 0.967)
114 surahs SEPARATE	Exp 2: 26% of surahs mix with poets (silhouette 0.469)	Exp 8: 0% of surahs mix with poets (silhouette 0.679)

The two rows tell different stories. Restricting to well-documented poets eliminated the overlap when surahs are treated separately (Exp 2 to Exp 8), but only partially reduced it — and did not eliminate it — when the Quran is treated as one combined entity (Exp 1 to Exp 7). The Quran's own cluster and its top similarity score barely moved between Exp 1 and Exp 7.




Figure 1. Both comparison axes side by side. Left: overlap persists. Right: overlap disappears.

Experiment 7 in Detail

Restricted to the 35 well-documented poets, the Quran (as one combined entity) still landed inside a defined cluster, now with 9 poets instead of 19:

Aws ibn Hajar, Al-Tufayl al-Ghanawi, Al-Mutalammis al-Dab'i, Al-Nabigha al-Dhubyani, Ta'abbata Sharran, Hatim al-Ta'i, Zayd al-Khayl al-Ta'i, Urwa ibn al-Ward, Qays ibn al-Haddadiyya.

Several of these are well-known, canonical pre-Islamic poets (Hatim al-Ta'i, Al-Nabigha al-Dhubyani) — not obscure figures. Their similarity to the Quran (0.967) is nearly as high as the original 0.976 from Experiment 1.




Figure 2. Experiment 7. Red star = Quran, sitting among a cluster of well-documented poets (dark blue), separate from two other poet clusters.

What Changed Between the Two Rows?

The most likely explanation is a difference in what's being compared, not a difference in poet quality. In the top row, the Quran is compressed into a SINGLE vector — an average across all 114 surahs' worth of variation. That heavy averaging could itself produce a more "generic" or centered vector, regardless of which poets it's compared against, simply because averaging over a lot of diverse text tends to smooth out distinguishing features. In the bottom row, each surah keeps its own individual signature, so the comparison is more like-for-like — and that's exactly the comparison where restricting to well-documented poets produced total separation.

In short: the sample-size explanation looks well-supported for surah-level comparisons, but doesn't fully explain the whole-Quran-as-one-entity result. That second result may be more about how averaging 114 surahs into one vector affects distinctiveness than about the poets on the other side of the comparison. This wasn't tested directly here and would need a further experiment (e.g. comparing the Quran's whole-text vector against similarly averaged multi-poem vectors from a few individual poets) to confirm.

Overall Verdict

Across all four experiments: pre-Islamic poetry and the Quran are NOT maximally distinct in every representation tested. A small, fairly consistent set of poets — including some well-known, canonical figures, not just obscure ones — keep landing close to the Quran's style profile when it's treated as a single unified text. That overlap fully disappears, however, once the Quran's surahs are compared individually against equally well-documented poets. The honest summary is: the result depends on how the comparison is set up, and that dependency is itself informative — it points to averaging effects mattering as much as, or more than, which poets are included.

What This Does and Does Not Prove

Across all eight experiments run in this project, the same limit applies: this measures writing-style similarity, not authorship. Every result here describes a pattern in how text embeddings sit relative to each other under one specific method. None of it identifies, confirms, or rules out who wrote what. That interpretive boundary hasn't moved across any version of this analysis.